

- 1) the CPU is our bottleneck, it can't make computations in the needed rate to reach FSB Target, the sequential nature of CPU execution limits performance.
- 2) FPGAs are ideal for accelerating image blur due to:
  - Parallelism – Process multiple pixels simultaneously.
  - Custom Hardware Implementation – Directly execute blur operations without CPU overhead.
  - Pipelining – Increase throughput by structuring computation in stages.
  - Optimized Memory Access – Reduce latency by efficiently handling data movement.
- 3) **Parallelism** means executing multiple operations simultaneously. Unlike CPUs, which process tasks sequentially:
  - FPGAs use **custom hardware circuits** to perform many computations in parallel.
  - **Pipelining** allows different parts of an operation to execute concurrently, like an assembly line.
  - Using Specialized Logic if needed, FPGAs are made up of configurable logic blocks that can be arranged to perform specific operations very efficiently that we can arrange and use if needed.
- 4)
  - Profiling and Analysis – Identify bottlenecks and analyze data dependencies.
  - Algorithm Transformation – Restructure code for parallel execution.
  - Hardware Design – Implement using Verilog or High-Level Synthesis (HLS) tools.
  - Verification and Testing – Validate functionality via simulation and real FPGA testing.
- 5)

FPGA Clock Speed is Typically in the 100 MHz – 500 MHz range while ,CPU Clock Speed is Usually in the GHz range (e.g., 3 GHz).

Despite lower clock speeds, **FPGAs outperform CPUs** because they:

  - Execute **many operations in parallel** instead of sequentially.
  - Use **custom circuits** for specific tasks, avoiding instruction decoding overhead.
  - Employ **pipelining**, improving efficiency despite lower frequency.

Basically, this is due to FPGA Parallelism.

6) The host CPU acts as the controller and is responsible for:

- Initializing the FPGA and allocating memory.
- Transferring data between CPU and FPGA memory.
- Launching kernels (hardware-accelerated functions on the FPGA).
- Handling sequential logic isn't efficient for the FPGA.
- Retrieving results and performing final processing.

7)

- Data Locality – Keep frequently used data on the FPGA as long as possible.
- Coalesced Memory Access – Optimize data layout to minimize unnecessary memory reads/writes.
- Overlapping Computation and Communication – Transfer data while computation is still running.
- Efficient Direct Memory Access (DMA) – Use fast data transfer methods instead of CPU-managed transfers.

8)

```
void printImage(uint8_t *image, int M, int N) {  
  
    int paddedN = (N % 4 == 0) ? N : N + (4 - N % 4);  
  
    uint8_t output[M + 1][paddedN];  
  
  
    // Copy image and pad last row and column  
    for (int i = 0; i < M; i++) {  
        for (int j = 0; j < N; j++) {  
            output[i][j] = image[i * N + j];  
        }  
        for (int j = N; j < paddedN; j++) {  
            output[i][j] = output[i][N - 1]; // Duplicate last column  
        }  
    }  
  
    for (int j = 0; j < paddedN; j++) {  
        output[M][j] = output[M - 1][j]; // Duplicate last row
```

```
}

// Print words (4 bytes at a time)
for (int i = 0; i < M + 1; i++) {
    for (int j = 0; j < paddedN; j += 4) {
        printf("{%d, %d, %d, %d}\n",
            output[i][j], output[i][j + 1],
            output[i][j + 2], output[i][j + 3]);
    }
}
}
```

9)

```
void printImage(uint8_t *image, int M, int N) {  
    int paddedN = (N % 4 == 0) ? N : N + (4 - N % 4);  
    uint8_t output[M + 1][paddedN];  
  
    // Copy image and pad last row and column  
    for (int i = 0; i < M; i++) {  
        for (int j = 0; j < N; j++) {  
            output[i][j] = image[i * N + j];  
        }  
        for (int j = N; j < paddedN; j++) {  
            output[i][j] = output[i][N - 1]; // Duplicate last column  
        }  
    }  
    for (int j = 0; j < paddedN; j++) {  
        output[M][j] = output[M - 1][j]; // Duplicate last row  
    }  
  
    // Print interleaved first and second outputs  
    for (int i = 0; i < M + 1; i++) {  
        for (int j = 0; j < paddedN; j += 4) {  
            printf("{%d, %d, %d, %d}\n", output[i][j], output[i][j + 1], output[i][j + 2], output[i][j + 3]);  
            if (i > 0) // Second printout starts from the second row  
                printf("{%d, %d, %d, %d}\n", output[i - 1][j], output[i - 1][j + 1], output[i - 1][j + 2],  
output[i - 1][j + 3]);  
        }  
    }  
}
```

