

1)

```
1 // Conceptual pseudocode
2 tic();
3 read_from(probe); // T1: Probe to CPU, Interference 1
4 toc();
5 T1 = time_elapsed();
6
7 tic();
8 write_to(memory); // T2: CPU to memory, Interference 2
9 toc();
10 T2 = time_elapsed();
11
12 tic();
13 read_from(memory); // T3: Memory to CPU, Interference 3
14 toc();
15 T3 = time_elapsed();
16
17 tic();
18 process_image(); // Image processing in CPU
19 toc();
20 process_time = time_elapsed();
21
22 tic();
23 write_to(monitor); // T4: CPU to monitor
24 toc();
25 T4 = time_elapsed();
26
27 print("T1: ", T1, " T2: ", T2, " T3: ", T3, " Process Time: ", process_time, " T4: ", T4);
28 |
```

2)

- Design level
- Algorithm and data structure level
- Source code level
- Build level
- Compile level
- Assembly level
- Run time (e.g., Just-in-Time compilation)

General order: Starting from higher levels (design, algorithms) to lower levels (source, build, compile, assembly).

Since refinement at higher levels often has a larger impact and is harder to change later and could result in a complete re-writing.

3)

- **CPU-bound:** Task that depend heavily on the CPU, with long CPU bursts and infrequent I/O (or any other) operations. Performance improves by increasing CPU speed.
- **I/O-bound:** Tasks depend on the input/output system, with frequent, short CPU work with long waits for I/O between each CPU burst. Performance improves by speeding up I/O operations.

4)

- Resolution: $X \times Y \times Y \times Y$ pixels per image, 1 byte per pixel
- Frame rate: $15 \leq FPS \leq 30$
- Latency: No more than 500 ms
- Size/weight: Comparable to a regular PC

5)

Optimizations often add complexity (e.g., special cases, tricks) and reduce readability, making debugging and updates more difficult.

Still, we will accept this trade-off, when achieving critical performance benchmarks or resolving bottlenecks in time-sensitive systems.

How to minimize the negative impact:

- Use modular and well-documented code.
- Isolate optimized code sections and add comments/refers to explain them.

6)

To identify a bottleneck in a complex system we can either use profiling tools to measure resource utilization or analyse critical sections of the code (hot spots) and identify resource constraints (e.g., CPU, memory, I/O).

To the primary performance limiter after identifying the bottleneck

- Experiment with changes to the suspected bottleneck and measure the performance impact.

- Validate those improvements in the bottleneck lead to overall system performance gains.
- Rule out other potential bottlenecks through systematic testing and analysis.